

KLE Society's  
KLE Technological University



# DISTRIBUTED AND CLOUD COMPUTING

Course Code: 20ECSC305

## CICD PIPELINE ON KUBERNETES USING AWS

### TEAM - 02

Shreeya Goggi	01fe19bcs045	140
Rashmi Kiragi	01fe19bcs057	150
Renuka Talwar	01fe19bcs068	159
Sahana Bhasme	01fe19bcs072	163

**COURSE TEACHER**  
**Priyadarshini Patil**

**SCHOOL OF COMPUTER SCIENCE & ENGINEERING**  
**KLE TECHNOLOGICAL UNIVERSITY**  
**HUBLI-580 031 (India)**  
**Academic year 2021-22**

# **Abstract**

Kubernetes, also known as K8s, is an open-source system for automating deployment, scaling, and management of containerized applications across multiple hosts. It groups containers that make up an application into logical units for easy management and discovery. Kubernetes builds upon a decade and a half of experience at Google running production workloads at scale using a system called Borg, combined with best-of-breed ideas and practices from the community. Kubernetes is hosted by the Cloud Native Computing Foundation. Amazon, Google, IBM, Microsoft, Oracle, Red Hat, SUSE and VMWare offer Kubernetes-based platforms or infrastructure as a service (IaaS) that deploy Kubernetes.

# **Introduction**

## **Why Kubernetes ?**

Kubernetes allows us to deploy cloud-native applications anywhere and manage them everywhere. As most modern software developers can attest, containers have provided us with dramatically more flexibility for running cloud-native applications on physical and virtual infrastructure. Containers package up the services comprising an application and make them portable across different compute environments, for both dev/test and production use. With containers, it's easy to quickly ramp application instances to match spikes in demand. And because containers draw on resources of the host OS, they are much lighter weight than virtual machines. This means containers make highly efficient use of the underlying server infrastructure.

## **How are Kubernetes different ?**

### **Traditional deployment era:**

Early on, organizations ran applications on physical servers. There was no way to define resource boundaries for applications in a physical server, and this caused resource allocation issues. For example, if multiple applications run on a physical server, there can

be instances where one application would take up most of the resources, and as a result, the other applications would underperform. A solution for this would be to run each application on a different physical server. But this did not scale as resources were underutilized, and it was expensive for organizations to maintain many physical servers.

### **Virtualized deployment era:**

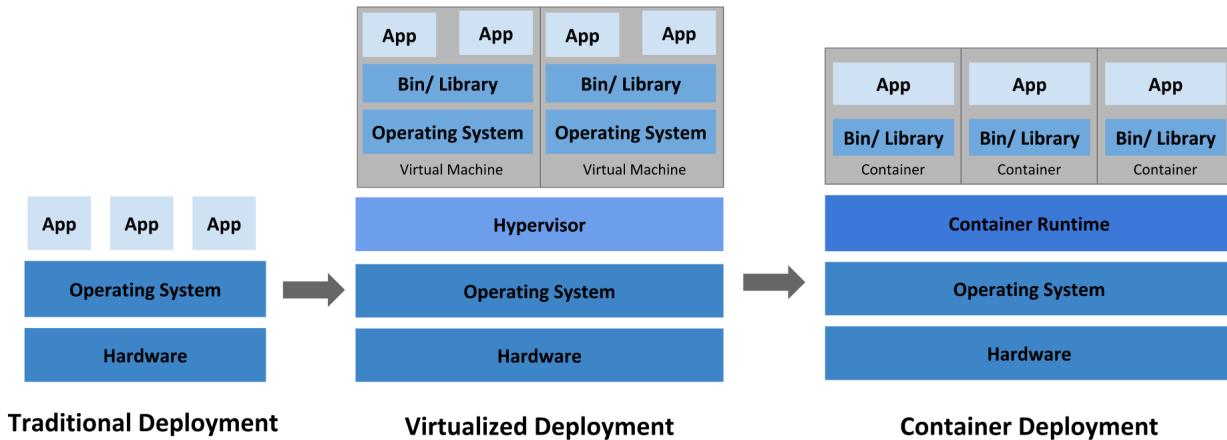
As a solution, virtualization was introduced. It allows you to run multiple Virtual Machines (VMs) on a single physical server's CPU. Virtualization allows applications to be isolated between VMs and provides a level of security as the information of one application cannot be freely accessed by another application.

Virtualization allows better utilization of resources in a physical server and allows better scalability because an application can be added or updated easily, reduces hardware costs, and much more. With virtualization you can present a set of physical resources as a cluster of disposable virtual machines.

Each VM is a full machine running all the components, including its own operating system, on top of the virtualized hardware.

### **Container deployment era:**

Containers are similar to VMs, but they have relaxed isolation properties to share the Operating System (OS) among the applications. Therefore, containers are considered lightweight. Similar to a VM, a container has its own filesystem, share of CPU, memory, process space, and more. As they are decoupled from the underlying infrastructure, they are portable across clouds and OS distributions.



## Benefits of Containers

- Agile application creation and deployment: increased ease and efficiency of container image creation compared to VM image use.
- Continuous development, integration, and deployment: provides for reliable and frequent container image build and deployment with quick and efficient rollbacks (due to image immutability).
- Dev and Ops separation of concerns: create application container images at build/release time rather than deployment time, thereby decoupling applications from infrastructure.
- Observability: not only surfaces OS-level information and metrics, but also application health and other signals.
- Environmental consistency across development, testing, and production: Runs the same on a laptop as it does in the cloud.
- Cloud and OS distribution portability: Runs on Ubuntu, RHEL, CoreOS, on-premises, on major public clouds, and anywhere else.
- Application-centric management: Raises the level of abstraction from running an OS on virtual hardware to running an application on an OS using logical resources.

- Loosely coupled, distributed, elastic, liberated micro-services: applications are broken into smaller, independent pieces and can be deployed and managed dynamically – not a monolithic stack running on one big single-purpose machine.
- Resource isolation: predictable application performance.
- Resource utilization: high efficiency and density.

## Gaps And Challenges

Enterprises are quickly discovering new Kubernetes use cases to increase agility and application performance. While Kubernetes provides more freedom to run applications across a variety of infrastructures, broader adoption means a proliferation of new clusters and increased complexity in how containers are run, monitored and managed. While managing one Kubernetes cluster is not trivial, trying to manage multiple Kubernetes clusters on hybrid or multiple clouds becomes exponentially more difficult. There are a number of common challenges organizations run into when attempting to efficiently rein in multiple clusters. Developers should keep the following in mind when developing a strategy to overcome these three common governance challenges in Kubernetes deployments:

### **Lack of Visibility and Management:**

As the number of clusters grows and spreads, managing and tracking their activity and growth becomes increasingly difficult. It is also harder and more time-consuming to troubleshoot any problems that may arise; if different software is involved, a single solution cannot be applied to each version. Lack of centralized governance and visibility into what's happening within the Kubernetes environment negatively impacts application availability and performance and, ultimately, the organization's bottom line.

Problems also arise if a cluster goes down unexpectedly—troubleshooting problems requires time and resources, and if there are dozens of potential software versions in use, managing all of them across the organization is increasingly difficult. Unlike planned downtime, when teams are generally able to understand what is needed to mitigate impacts for customers and projects, unplanned downtime makes this preparation impossible. Operators must be able to consistently administer, manage and obtain insights about their infrastructure.

### **Operational Complexity and Overhead:**

Spreading numerous Kubernetes clusters across different business units creates challenges in user identity tracking—especially if users onboard, offboard or change teams. Operators lose the flexibility to define user roles, responsibilities and privileges to ensure the right people are performing the right tasks within the environment. They also face difficulties identifying role violations, assessing governance risks and performing compliance checks. When more time is spent chasing issues or putting out fires, there is less time for efficient operations.

### **Empowering Both Developers and Operators:**

As developers work to implement Kubernetes, there must be a balance between the developer's independence and the operator's ability to easily manage the policies and procedures necessary to maintain the overall system. While the autonomy of developers is crucial, this freedom shouldn't come at the expense of the environment's security. Multiple solutions across multiple clusters often lead to more opportunities for attacks to slip through the cracks unnoticed. The alternative approach, where there is more consistency for a few select solutions, allows standardization across organizational clusters. The challenge comes when addressing the fine line between a developer's ability to innovate and an operator's ability to maintain the necessary procedures and governance.

## **Problem Statement**

### **To Build a CI/CD Pipeline for Microservices on Kubernetes.**

#### **Objectives**

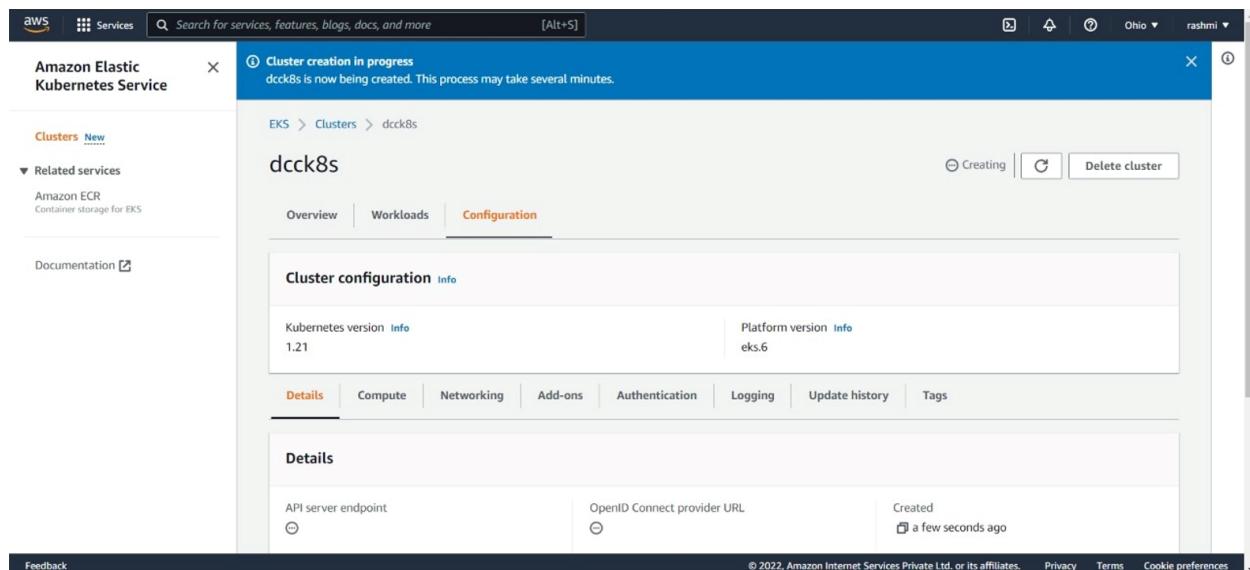
- To build and deploy services independently.
- Code changes that pass the CI process are automatically deployed to a production-like environment.
- A new version of a service can be deployed side by side with the previous version.

## Tools

- AWS EC2 instance
- AWS EKS
- Jenkins
- Codecommit pipeline
- AWS ECS
- Github repository
- Kubectl
- Eksctl
- AWS cli

## Program implementation plan

To set up a Kubernetes cluster consisting of one Master node and one or more worker nodes.



# RSA public key generator

```
apiVersion: v1
clusters:
- cluster:
    certificate-authority-data:
        LS0tLS1CRUdJTIBDRVJUSUZJQ0FURS0tLS0tCk1JSUM1ekNDQWMrZ0F3SUJBZ01CQURBTkJna
        3Foa2lHOXcwQkFRc0ZBREFWTVJNd0VRWURWUVFERXdwcmRXSmwKY201bGRHVnpNQjRyRFRJeU
        1EUXhNekE0TWpVME1Gb1hEVE15TURRe1EQTRNa1UwTUZvd0ZURVRNqKVHQTfVRQpBeE1LYTN
        WaVpY8nVaWFJsY3pDQ0FTSXdEUV1KS29aSwH2Y05BUUVCQ1FBRGdnRVBBRENDQVFvQ2dnRUJB
        S2g5Cn1DV1U5WGE0ZmhZUk1MVHdzbk56R1VXNjlialpaY0tsb3JKZmxUMzFIdHF4eFRTTFhtM
        DdNWUs0UTNWl3YydlcKak0vS0ZKWEJsN1plUVhvNjAzTzdpbk850GQzMKVQSkZGcUz1N1VqT1
        1VSXVDcFV1c1RYTHJUeGd0bWZTVDRmbgpIbXEvrHJpVDIxemFzc3I1YkQ3MzU0NkZPL2w3OXN
        LUklWdDBTVnhNYk1LTXF1V2FFMjMvYkhLa3JKNXN2MLFPC1A2NTRnVGg4UzFsaDhHY2h1ZXN4
        UEpiRTY0aGdHcXUvcWJFb11ZnF4UjN3Y0N3OHN2amhBSkJKck5jZ1E5VGkKa2xLNEVJSnBib
        jdJUEVNY1Uch1FUW1DMVNQV1lQRDJxkEpCd2RCMnM3L2pLL2tRV1E3S1RochHMxTk5KU1pwQw
        plUU9tenUvcHhRWVJVb0gxq3YwQ0F3RUFBYU5DTUVBd0RnWURWUjBqqVFIL0JBUIURBz0trTUE
        4ROexVWRFd0VCCI93UUZNUQ1CQWY4d0hRURWUjBPQkJZRUZZCZmozt19MOFZsbF1OZGpWeE1Q
        djRJcE05Z0tNQTBHQ1NxR1nJYjMKRFFFQkN3VUFBNE1CQVFCVmld1VmJDT0hXUUxmZG11M0pYM
        FVnL3ZZckRKVVowT3BHWFM2SXVZT01ZkwccndtegpROWY3TFFJb0VDV1RHeW52RERqOTryT0
        xZMVFGUH1ydENiZHM0WXZpdHpCY2ZPNFBTeWozi9Edip1WgtWUFLockVITnB3UXV45nFrRFh
        HNjhQcXp2WUQ3MURQUUJ5NXYzQ254U0dqTUxRc1RpVTNINFJMREVNWdXS1VMNWE1RFMKaWgr
        OGx0Y2J3V3BjdUFje1RONWxWKz12q2hrWDhNaDYvUzQzS316ald1U0h1K1Nsb0x1MUZ1NWt1U
        UI5WUozYwcoyTDE3N3d20E9qMEE2c1dVOUCxcU4rZU1IwMznUnFNUjh6V0ZmNFpTWTZ1L3dZYz
        RqV012WndqSmdaM1oxQ1JZCmlwZ05RQkVsMG4rWXpOVjRHRjk2RGxVNmc4T2hHNk820ENveQo
        tLS0tLUVORCBDRVJUSUZJQ0FURS0tLS0tCg==
            server: https://6491FBC870261581C2DB82F94B9F5E72.gr7.us-east-
2.eks.amazonaws.com
    name: arn:aws:eks:us-east-2:803947141637:cluster/dcck8s
contexts:
- context:
    cluster: arn:aws:eks:us-east-2:803947141637:cluster/dcck8s
    user: arn:aws:eks:us-east-2:803947141637:cluster/dcck8s
    name: arn:aws:eks:us-east-2:803947141637:cluster/dcck8s
current-context: arn:aws:eks:us-east-2:803947141637:cluster/dcck8s
... . . .
```

# Creation of worker nodes

```
Command Prompt
:\>aws configure
AWS Access Key ID [*****NUKA]: AKTA3WLX33IC43GTNQWT
AWS Secret Access Key [None]: N6P6+75hKkn/kskfNqGMCH00S6thc/2uTb8o2p9J
default region name [None]: us-east-2
default output format [None]: json

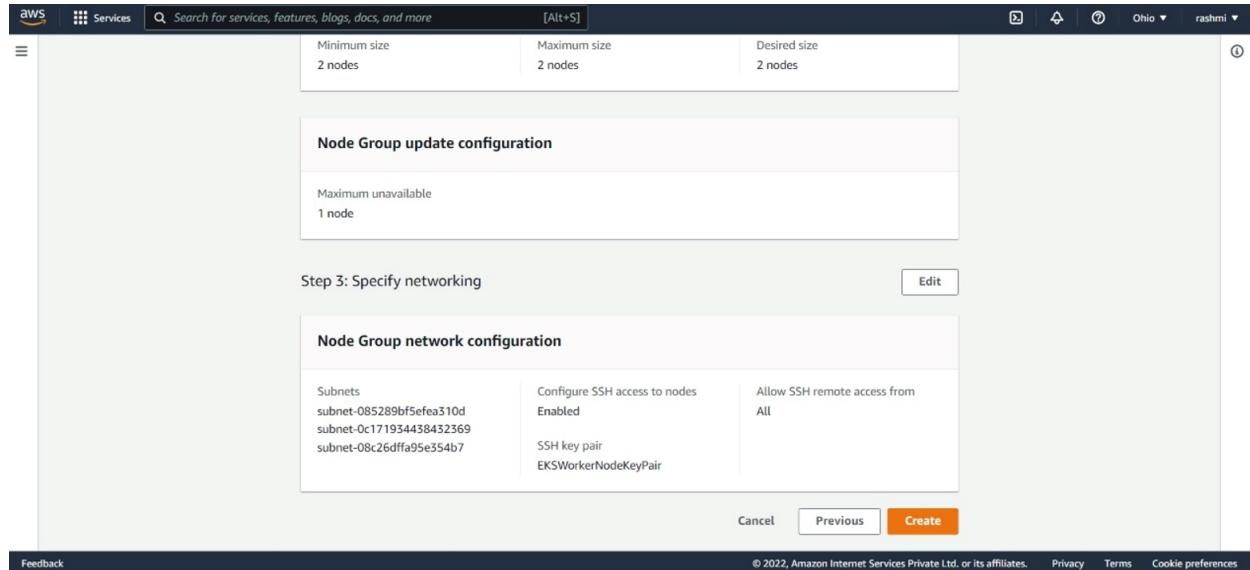
:\>aws eks --region us-east-2 describe-cluster --name dcck8s --query cluster.status
ACTIVE"

:\>aws eks --region us-east-2 update-kubeconfig --name dcck8s
added new context arn:aws:eks:us-east-2:803947141637:cluster/dcck8s to C:\Users\DELL\.kube\config

:\>kubectl get svc
NAME          TYPE      CLUSTER-IP   EXTERNAL-IP   PORT(S)    AGE
kubernetes   ClusterIP  10.100.0.1  <none>        443/TCP   87m

:\>kubectl get nodes --watch
NAME                               STATUS   ROLES   AGE     VERSION
p-172-31-0-230.us-east-2.compute.internal   Ready   <none>  3m11s  v1.21.5-eks-9017834
p-172-31-26-23.us-east-2.compute.internal   Ready   <none>  3m24s  v1.21.5-eks-9017834
p-172-31-26-23.us-east-2.compute.internal   Ready   <none>  5m31s  v1.21.5-eks-9017834
p-172-31-0-230.us-east-2.compute.internal   Ready   <none>  5m32s  v1.21.5-eks-9017834

:\>kubectl get nodes
NAME                               STATUS   ROLES   AGE     VERSION
p-172-31-0-230.us-east-2.compute.internal   Ready   <none>  6m44s  v1.21.5-eks-9017834
p-172-31-26-23.us-east-2.compute.internal   Ready   <none>  6m57s  v1.21.5-eks-9017834
```



# Configuring AWS cli

```
Command Prompt

D:\>aws eks --region us-east-2 describe-cluster --name dcck8s --query cluster.status
Unable to locate credentials. You can configure credentials by running "aws configure".

D:\>aws configure
AWS Access Key ID [None]: RASHMIKIRAGI14RENUKA
AWS Secret Access Key [None]:
Default region name [None]:
Default output format [None]:

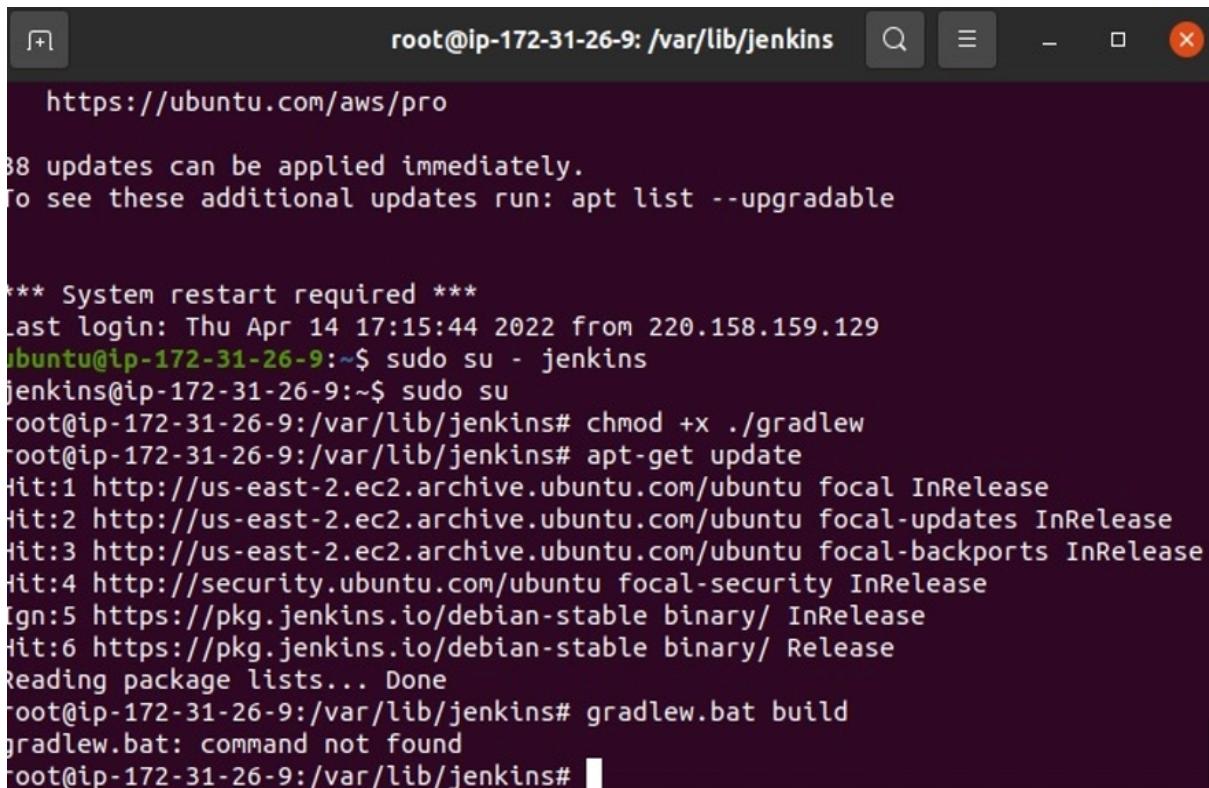
D:\>aws configure
AWS Access Key ID [*****NUKA]: AKIA3WLX33IC43GTNQWI
AWS Secret Access Key [None]: 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
Default region name [None]: us-east-2
Default output format [None]: json

D:\>aws eks --region us-east-2 describe-cluster --name dcck8s --query cluster.status
"ACTIVE"

D:\>aws eks --region us-east-2 update-kubeconfig --name dcck8s
Added new context arn:aws:eks:us-east-2:803947141637:cluster/dcck8s to C:\Users\DELL\.kube\config

D:\>
```

# Creating CI/CD pipeline using jenkins



A terminal window titled "root@ip-172-31-26-9: /var/lib/jenkins". The window contains the following command-line session:

```
root@ip-172-31-26-9: /var/lib/jenkins
https://ubuntu.com/aws/pro
38 updates can be applied immediately.
To see these additional updates run: apt list --upgradable

*** System restart required ***
Last login: Thu Apr 14 17:15:44 2022 from 220.158.159.129
ubuntu@ip-172-31-26-9:~$ sudo su - jenkins
jenkins@ip-172-31-26-9:~$ sudo su
root@ip-172-31-26-9:/var/lib/jenkins# chmod +x ./gradlew
root@ip-172-31-26-9:/var/lib/jenkins# apt-get update
Hit:1 http://us-east-2.ec2.archive.ubuntu.com/ubuntu focal InRelease
Hit:2 http://us-east-2.ec2.archive.ubuntu.com/ubuntu focal-updates InRelease
Hit:3 http://us-east-2.ec2.archive.ubuntu.com/ubuntu focal-backports InRelease
Hit:4 http://security.ubuntu.com/ubuntu focal-security InRelease
Ign:5 https://pkg.jenkins.io/debian-stable binary/ InRelease
Hit:6 https://pkg.jenkins.io/debian-stable binary/ Release
Reading package lists... Done
root@ip-172-31-26-9:/var/lib/jenkins# gradlew.bat build
gradlew.bat: command not found
root@ip-172-31-26-9:/var/lib/jenkins#
```

# Pushing the application to Github repository and compiling using gradle and building

```
$ ls
Hp@rashmik MINGW64 ~/node-api (master)
$ git status
On branch master
No commits yet
nothing to commit (create/copy files and use "git add" to track)

Hp@rashmik MINGW64 ~/node-api (master)
$ cd ..

Hp@rashmik MINGW64 ~
$ git clone ssh://APKA3WLX33ICU2IIF7ES@git-codecommit.us-east-2.amazonaws.com/v1/repos/node-api
fatal: destination path 'node-api' already exists and is not an empty directory.

Hp@rashmik MINGW64 ~
$ git clone ssh://APKA3WLX33ICU2IIF7ES@git-codecommit.us-east-2.amazonaws.com/v1/repos/node-api
Cloning into 'node-api'...
remote: Counting objects: 35, done.
Receiving objects: 100% (35/35), 32.45 KiB | 2.70 MiB/s, done.

Hp@rashmik MINGW64 ~
$ cd node-api

Hp@rashmik MINGW64 ~/node-api (main)
$ code .

Hp@rashmik MINGW64 ~/node-api (main)
$ git add buildspec.yml

Hp@rashmik MINGW64 ~/node-api (main)
$ git commit -m "Add buildspec"
[main 5d9c05e] Add buildspec
 1 file changed, 16 insertions(+)
 create mode 100644 buildspec.yml

Hp@rashmik MINGW64 ~/node-api (main)
$ git push remote main
fatal: 'remote' does not appear to be a git repository
fatal: Could not read from remote repository.

Please make sure you have the correct access rights
and the repository exists.

Hp@rashmik MINGW64 ~/node-api (main)
$ git push origin main
Enumerating objects: 4, done.
Counting objects: 100% (4/4), done.
Delta compression using up to 8 threads
Compressing objects: 100% (3/3), done.
Writing objects: 100% (3/3), 392 bytes | 196.00 KiB/s, done.
Total 3 (delta 1), reused 0 (delta 0), pack-reused 0
To ssh://git-codecommit.us-east-2.amazonaws.com/v1/repos/node-api
 13c9731..5d9c05e main -> main

Hp@rashmik MINGW64 ~/node-api (main)
$
```

The screenshot shows the AWS CodeCommit interface. The top navigation bar includes the AWS logo, Services, a search bar, and user information (Rashmi @ 8039-4714-1637). Below the navigation is a breadcrumb trail: Developer Tools > CodeCommit > Repositories > node-api. The main content area displays the 'node-api' repository with a list of files:

- test
- .dockercfg
- .gitignore
- .nvmrc
- app.js
- buildspec.yml
- Dockerfile
- package-lock.json
- package.json
- README.md
- run.js

On the right side of the file list, there are buttons for 'Notify', 'main', 'Create pull request', 'Clone URL', and 'Add file'. The left sidebar contains navigation links for Source, Code, Artifacts, Build, Deploy, Pipeline, and Settings, along with links for Go to resource and Feedback.

The screenshot shows a terminal window with the following output:

```
root@ip-172-31-26-9: /var/lib/jenkins
+ ueue.xml
+ ueue.xml.bak
+ ecret.key
+ ecret.key.not-so-secret
+ ecrets
+ pdates
+ serContent
+ sers
+ orkflow-libs
+ orkspace
oot@ip-172-31-26-9:/var/lib/jenkins# unzip gradle-6.4.1-bin.zip
archive: gradle-6.4.1-bin.zip
  creating: gradle-6.4.1/
  inflating: gradle-6.4.1/README
  inflating: gradle-6.4.1/LICENSE
  inflating: gradle-6.4.1/NOTICE
  creating: gradle-6.4.1/init.d/
  inflating: gradle-6.4.1/init.d/readme.txt
  creating: gradle-6.4.1/bin/
  inflating: gradle-6.4.1/bin/gradle
  inflating: gradle-6.4.1/bin/gradle.bat
  creating: gradle-6.4.1/lib/
```

```
root@ip-172-31-26-9: /var/lib/jenkins
.....
Unzipping /root/.gradle/wrapper/dists/gradle-6.4.1-all/13imxtezgn9nwzqt8rgtkunh1/gradle-6.4.1-all.zip to /root/.gradle/wrapper/dists/gradle-6.4.1-all/13imxtezgn9nwzqt8rgtkunh1
Set executable permissions for: /root/.gradle/wrapper/dists/gradle-6.4.1-all/13imxtezgn9nwzqt8rgtkunh1/gradle-6.4.1/bin/gradle

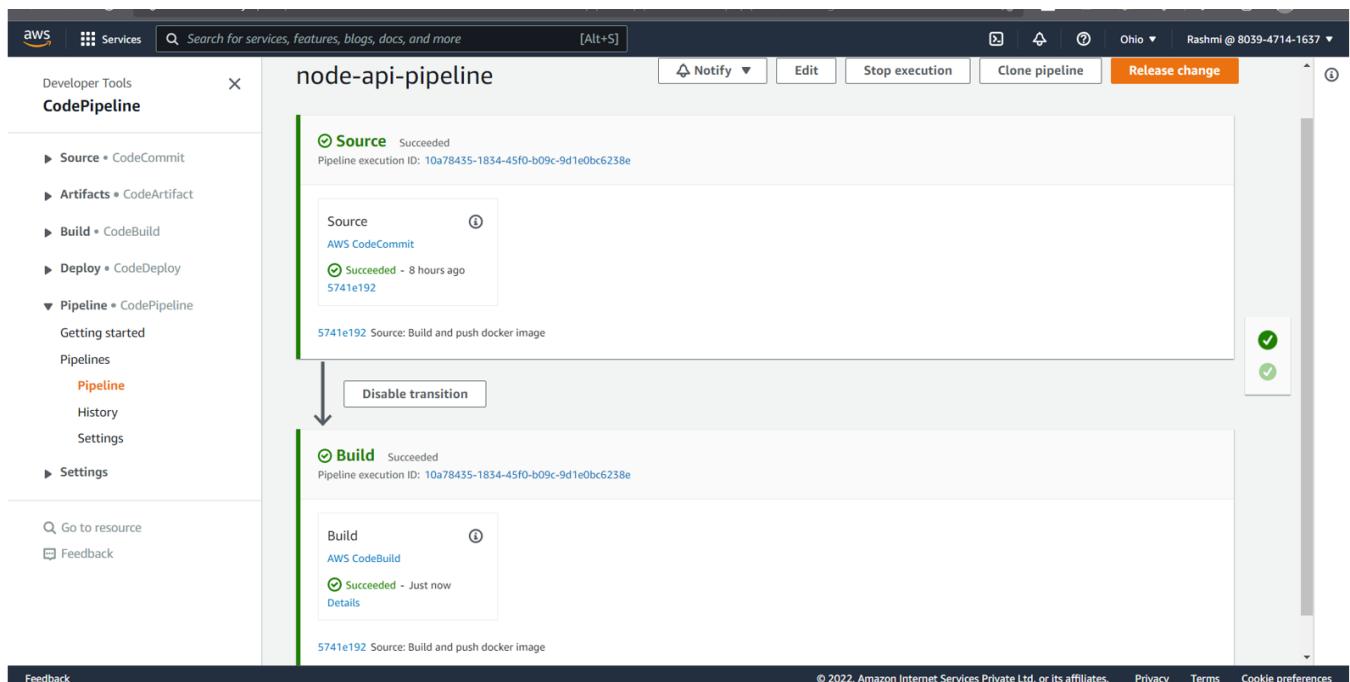
Welcome to Gradle 6.4.1!

Here are the highlights of this release:
- Support for building, testing and running Java Modules
- Precompiled script plugins for Groovy DSL
- Single dependency lock file per project

For more details see https://docs.gradle.org/6.4.1/release-notes.html

Starting a Gradle Daemon (subsequent builds will be faster)

Deprecated Gradle features were used in this build, making it incompatible with Gradle 7.0.
Use '--warning-mode all' to show the individual deprecation warnings.
See https://docs.gradle.org/6.4.1/userguide/command\_line\_interface.html#sec:command-line-deprecation-warnings
```



**Build logs**

⌚ Succeeded      Start time: 3 hours ago      Current phase: COMPLETED

```

185  c2ba5b92b5b0: Waiting
186  0cd41aa8001f: Waiting
187  7a6d0f54488f: Waiting
188  6ef8823b489f: Waiting
189  9578c16f3f7c: Waiting
190  3141322c5cdb: Waiting
191  73dbf42c0a44: Pushed
192  4bac553c3495: Pushed
193  5d98f342bd35: Pushed
194  b0d4df5d1936b: Pushed
195  b9ced5dcda4c: Pushed
196  c26a5692b560: Pushed
197  1307c965765b: Pushed
198  6ef8823b489f: Pushed
199  9578c16f3f7c: Pushed
200  9b5ede322cf7: Pushed
201  3141322c5cdb: Pushed
202  7a6d0f54488f: Pushed
203  0cd41aa8001f: Pushed
204  latest: digest: sha256:0637d9fddbe9bd055d761452b01109d9edc25c4299472241197a7ccebb58562b size: 3051
205
206 [Container] 2022/04/20 05:14:05 Phase complete: POST_BUILD State: SUCCEEDED
207 [Container] 2022/04/20 05:14:05 Phase context status code: Message:
208

```

# Configuring a codecommit pipeline service to deploy

The screenshot shows the AWS CodePipeline console. On the left, there's a sidebar with navigation links for Developer Tools, CodePipeline, Source, Artifacts, Build, Deploy, Pipeline, and Pipeline settings. The main area displays a pipeline named "node-api-pipeline". The pipeline has one stage, "Source", which is listed as "Succeeded". Below the pipeline table, there are links for "Go to resource" and "Feedback". At the bottom, there are footer links for "Feedback", "© 2022, Amazon Internet Services Private Ltd. or its affiliates.", "Privacy", "Terms", and "Cookie preferences".

Name	Most recent execution	Latest source revisions	Last executed
node-api-pipeline	⌚ Succeeded	Source - 5741e192: Build and push docker image	3 hours ago

# Deployment

The screenshot shows the AWS CodeCommit interface. The left sidebar is titled "CodeCommit" and includes sections for Source, Artifacts, Build, Deploy, Pipeline, and Settings. The main content area displays the "node-api" repository. At the top, there are buttons for Notify (dropdown), main (dropdown), Create pull request, and Clone URL. Below these are buttons for Add file (dropdown) and Add file. The repository listing shows the following files:

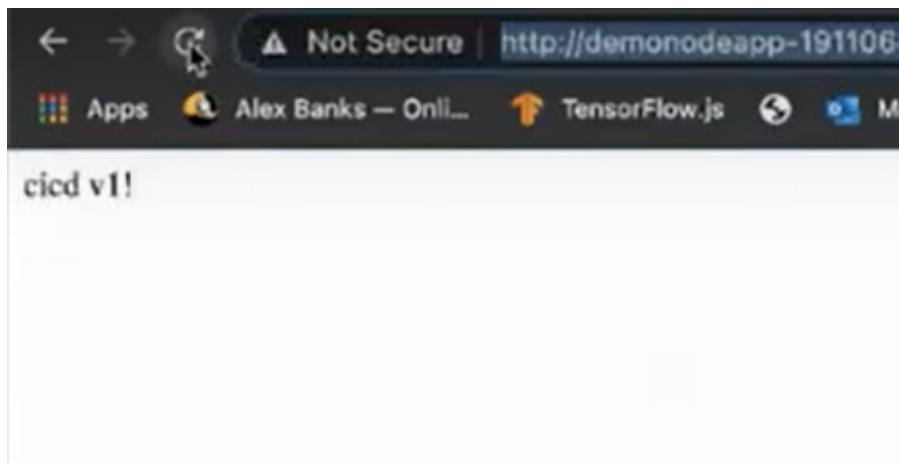
- test
- .dockerignore
- .gitignore
- .nvmrc
- app.js
- buildspec.yml
- Dockerfile
- package-lock.json
- package.json
- README.md
- run.js

At the bottom right of the main area, there are links for "© 2022, Amazon Internet Services Private Ltd. or its affiliates.", "Privacy", "Terms", and "Cookie preferences".

The screenshot shows the AWS Elastic Container Service (ECS) Cluster overview page for the "node-api" cluster. The top navigation bar includes the AWS logo, Services, a search bar, and user information (Ohio, Rashmi @ 8039-4714-1637). A green banner at the top states "node-api has been deployed successfully." The main content area shows the "Cluster overview" section with tabs for Services, Tasks, Infrastructure, Metrics, and Tags. The Services tab is selected, showing a table with columns for Service name, ARN, Status, Deployments and tasks, Task definition, Revision, and Launch type. The table currently displays one row: "arn:aws:ecs:us-east-2:803947141637:cluster/node-api". At the bottom of the page, there are links for "Feedback", "© 2022, Amazon Internet Services Private Ltd. or its affiliates.", "Privacy", "Terms", and "Cookie preferences".

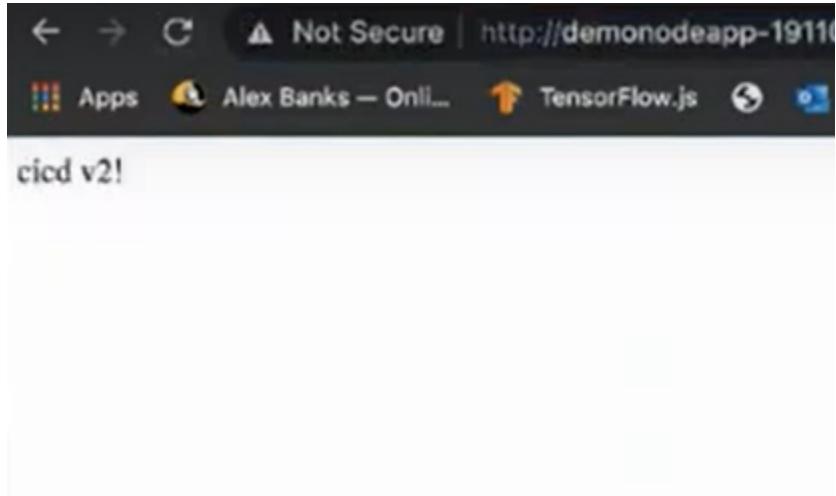
# Deployed application

Version v1



Version v2

Any changes in the code can be immediately seen in the web page due to the CI/CD pipeline.



# References

[\(59\) Kubernetes Tutorial for Beginners \[FULL COURSE in 4 Hours\] - YouTube](#)

[Getting started with Amazon EKS - Amazon EKS](#)

[What is IAM? - AWS Identity and Access Management \(amazon.com\)](#)

[AWS CodePipeline Documentation \(amazon.com\)](#)

[AWS CodeCommit Documentation \(amazon.com\)](#)

[\(59\) Build a CI/CD pipeline in AWS for container applications - YouTube](#)

[\(59\) CI/CD With AWS ECS + CodePipeline + CodeCommit + CodeBuild | DevOps With AWS Ep7 - YouTube](#)

[\(59\) AWS EKS CI/CD: Cluster Set-up, App containerization & Deployment, CodePipeline| DevOps With AWS Ep10 - YouTube](#)

[\(59\) AWS EKS CI/CD: Cluster Set-up, App containerization & Deployment, CodePipeline| DevOps With AWS Ep10 - YouTube](#)

[awscicd/app at main · abohmeed/awscicd \(github.com\)](#)